Experimento 2

[1]:
```
import torch
import torch.nn as nn
import torchvision
import torchvision.transforms as transform
import torch.nn.functional as F

import matplotlib.pyplot as plt
import numpy as np
import math

device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
#device='cpu'
print(device)

batch_size = 32

img_transform = transform.Compose([transform.ToTensor(), transform.Normalize((0.
 ↪5,),(0.5,))])

data = torchvision.datasets.ImageFolder("C:/Users/matia/Desktop/AI_UTEC/
 ↪proyecto5/fotos",transform=img_transform)

print(len(data))


train_set,test_set=torch.utils.data.random_split(data,[14815,6350],␣
 ↪generator=torch.Generator().manual_seed(0))
val_set,test_set=torch.utils.data.random_split(test_set,[4233,2117],␣
 ↪generator=torch.Generator().manual_seed(0))

img, _ = train_set[0]
print(img.shape)

train_loader = torch.utils.data.DataLoader(dataset=train_set,␣
 ↪batch_size=batch_size, shuffle=True)
test_loader = torch.utils.data.DataLoader(dataset=test_set,␣
 ↪batch_size=batch_size, shuffle=False)
```

```
val_loader = torch.utils.data.DataLoader(dataset=val_set, batch_size=batch_size,␣
 ↪shuffle=False)
```

```
cuda:0
21165
torch.Size([3, 299, 299])
```

[2]:
```
def show_img(img):
    plt.imshow(img.numpy()[0], cmap='gray')
```

[3]:
```
print(len(train_set))
print(len(test_set))
print(len(val_set))
```

```
14815
2117
4233
```

[4]:
```
img, label = train_set[999]
print(label)
show_img(img)
```

```
3
```

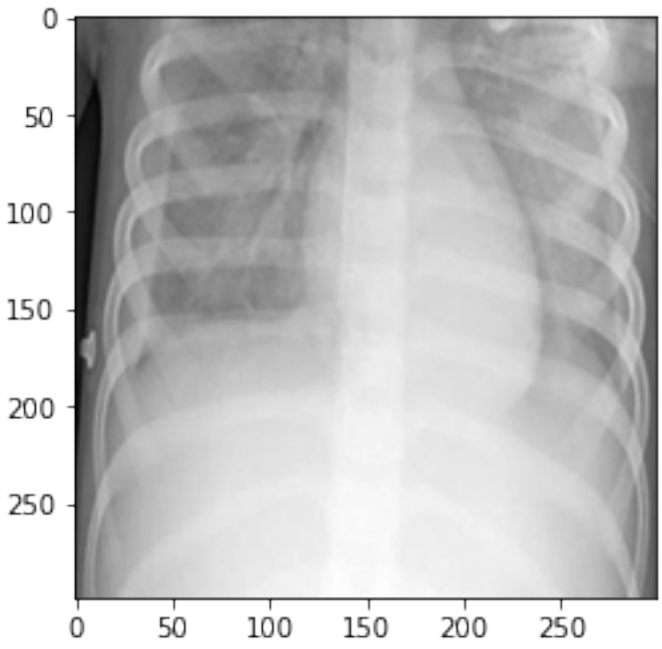

[5]:
```
#hyperparametros
num_classes = 4
learning_rate =  0.01
```

```
num_epochs = 20

class CNN(nn.Module):
    def __init__(self, num_classes=4):
        super(CNN, self).__init__()
        self.layer1 = nn.Sequential(
            nn.Conv2d(in_channels=3, out_channels=16, kernel_size=10,␣
 ↪stride=1,padding=0),
            nn.ReLU())
        self.layer2 = nn.Sequential(
            nn.Conv2d(in_channels=16,out_channels=32, kernel_size=6, stride=1,␣
 ↪padding=2),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2))
        self.layer3 = nn.Sequential(
            nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3, stride=1,␣
 ↪padding=2),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2))
        self.layer4 = nn.Sequential(
            nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3, stride=1,␣
 ↪padding=2),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2))
        self.fc = nn.Linear(37*37*128, num_classes)

    def forward(self, x):
        out = self.layer1(x)
        out = self.layer2(out)
        out = self.layer3(out)
        out = self.layer4(out)
        out = out.reshape(out.size(0), -1)
        out = self.fc(out)
        return out
```

[6]:
```
model         = CNN(num_classes).to(device)
loss_fn       = nn.CrossEntropyLoss()
optimizer     = torch.optim.Adam(model.parameters(), lr = learning_rate)
#loss_train    = train(model, optimizer, loss_fn, num_epochs)
#test(model)

print([ e.shape  for e in model.fc.parameters()])

model.fc.weight
```

```
[torch.Size([4, 175232]), torch.Size([4])]
```

[6]:
```
Parameter containing:
tensor([[ 1.9060e-03,  3.4415e-04,  5.2328e-05,  ...,  1.2915e-03,
          6.0591e-04, -6.3812e-04],
        [ 1.5982e-03, -1.6400e-03, -6.5031e-04,  ...,  1.3073e-03,
         -2.1961e-05, -2.3485e-03],
        [-1.5196e-04,  1.0452e-03, -5.0353e-04,  ...,  1.4927e-03,
          6.5300e-04,  2.0211e-03],
        [-1.9493e-03, -1.5882e-04, -2.2413e-03,  ..., -6.4911e-04,
         -1.4299e-03, -1.0236e-03]], device='cuda:0', requires_grad=True)
```

[7]:
```
def train(model, optimizer, loos_fn, num_epochs):
    loss_vals = []
    running_loss =0.0
    # train the model
    total_step = len(train_loader)

    list_loss= []
    list_time = []
    j=0

    for epoch in range(num_epochs):
        for i, (images, labels) in enumerate(train_loader):
            images = images.to(device)
            labels = labels.to(device)
            # forward
            output = model(images)
            loss   = loss_fn(output, labels)
            # change the params
            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

            list_loss.append(loss.item())
            list_time.append(j)
            j+=1
            #print(i,end=",")
            if (i+1) % 100 == 0:
                #print()
                print ('Epoch [{}/{}], Step [{}/{}], Loss: {:.4f}' .
 ↪format(epoch+1, num_epochs, i+1, total_step, loss.item()))

    print('Finished Training Trainset')
    return list_loss,list_time
```

[8]:
```
list_loss,list_time=train(model,optimizer,loss_fn,num_epochs)

plt.plot(list_time,list_loss,color="green", label="Accuracy")
```

```
plt.legend()
plt.xlabel("Time")
plt.ylabel("Accuracy")
plt.show()
```

```
Epoch [1/20], Step [100/463], Loss: 1.2170
Epoch [1/20], Step [200/463], Loss: 1.1619
Epoch [1/20], Step [300/463], Loss: 1.2960
Epoch [1/20], Step [400/463], Loss: 1.1681
Epoch [2/20], Step [100/463], Loss: 1.0581
Epoch [2/20], Step [200/463], Loss: 1.0721
Epoch [2/20], Step [300/463], Loss: 1.3485
Epoch [2/20], Step [400/463], Loss: 1.0444
Epoch [3/20], Step [100/463], Loss: 1.1047
Epoch [3/20], Step [200/463], Loss: 1.0933
Epoch [3/20], Step [300/463], Loss: 1.0818
Epoch [3/20], Step [400/463], Loss: 1.3065
Epoch [4/20], Step [100/463], Loss: 1.1350
Epoch [4/20], Step [200/463], Loss: 1.1389
Epoch [4/20], Step [300/463], Loss: 1.0255
Epoch [4/20], Step [400/463], Loss: 1.2788
Epoch [5/20], Step [100/463], Loss: 1.1458
Epoch [5/20], Step [200/463], Loss: 1.1862
Epoch [5/20], Step [300/463], Loss: 1.1572
Epoch [5/20], Step [400/463], Loss: 1.1031
Epoch [6/20], Step [100/463], Loss: 1.1641
Epoch [6/20], Step [200/463], Loss: 1.2140
Epoch [6/20], Step [300/463], Loss: 1.1591
Epoch [6/20], Step [400/463], Loss: 1.2148
Epoch [7/20], Step [100/463], Loss: 1.0920
Epoch [7/20], Step [200/463], Loss: 1.2884
Epoch [7/20], Step [300/463], Loss: 1.1501
Epoch [7/20], Step [400/463], Loss: 1.2383
Epoch [8/20], Step [100/463], Loss: 1.1200
Epoch [8/20], Step [200/463], Loss: 1.1371
Epoch [8/20], Step [300/463], Loss: 1.1676
Epoch [8/20], Step [400/463], Loss: 1.1040
Epoch [9/20], Step [100/463], Loss: 1.3566
Epoch [9/20], Step [200/463], Loss: 1.1848
Epoch [9/20], Step [300/463], Loss: 1.1364
Epoch [9/20], Step [400/463], Loss: 1.2150
Epoch [10/20], Step [100/463], Loss: 0.9769
Epoch [10/20], Step [200/463], Loss: 1.2156
Epoch [10/20], Step [300/463], Loss: 1.2124
Epoch [10/20], Step [400/463], Loss: 1.1565
Epoch [11/20], Step [100/463], Loss: 1.1831
Epoch [11/20], Step [200/463], Loss: 1.2134
Epoch [11/20], Step [300/463], Loss: 1.1297
```

```
Epoch [11/20], Step [400/463], Loss: 1.1836
Epoch [12/20], Step [100/463], Loss: 1.0867
Epoch [12/20], Step [200/463], Loss: 1.2105
Epoch [12/20], Step [300/463], Loss: 1.2926
Epoch [12/20], Step [400/463], Loss: 1.2141
Epoch [13/20], Step [100/463], Loss: 1.1525
Epoch [13/20], Step [200/463], Loss: 1.0914
Epoch [13/20], Step [300/463], Loss: 1.2012
Epoch [13/20], Step [400/463], Loss: 1.1652
Epoch [14/20], Step [100/463], Loss: 1.3184
Epoch [14/20], Step [200/463], Loss: 1.1367
Epoch [14/20], Step [300/463], Loss: 1.4085
Epoch [14/20], Step [400/463], Loss: 1.1453
Epoch [15/20], Step [100/463], Loss: 1.1277
Epoch [15/20], Step [200/463], Loss: 1.2675
Epoch [15/20], Step [300/463], Loss: 1.1358
Epoch [15/20], Step [400/463], Loss: 1.2266
Epoch [16/20], Step [100/463], Loss: 1.1988
Epoch [16/20], Step [200/463], Loss: 1.3594
Epoch [16/20], Step [300/463], Loss: 1.1721
Epoch [16/20], Step [400/463], Loss: 1.3141
Epoch [17/20], Step [100/463], Loss: 1.1503
Epoch [17/20], Step [200/463], Loss: 1.0854
Epoch [17/20], Step [300/463], Loss: 1.1865
Epoch [17/20], Step [400/463], Loss: 1.2475
Epoch [18/20], Step [100/463], Loss: 1.2370
Epoch [18/20], Step [200/463], Loss: 1.1041
Epoch [18/20], Step [300/463], Loss: 1.1687
Epoch [18/20], Step [400/463], Loss: 1.2607
Epoch [19/20], Step [100/463], Loss: 1.2816
Epoch [19/20], Step [200/463], Loss: 1.1843
Epoch [19/20], Step [300/463], Loss: 1.2922
Epoch [19/20], Step [400/463], Loss: 1.0686
Epoch [20/20], Step [100/463], Loss: 1.0100
Epoch [20/20], Step [200/463], Loss: 1.4263
Epoch [20/20], Step [300/463], Loss: 1.2155
Epoch [20/20], Step [400/463], Loss: 1.0680
Finished Training Trainset
```

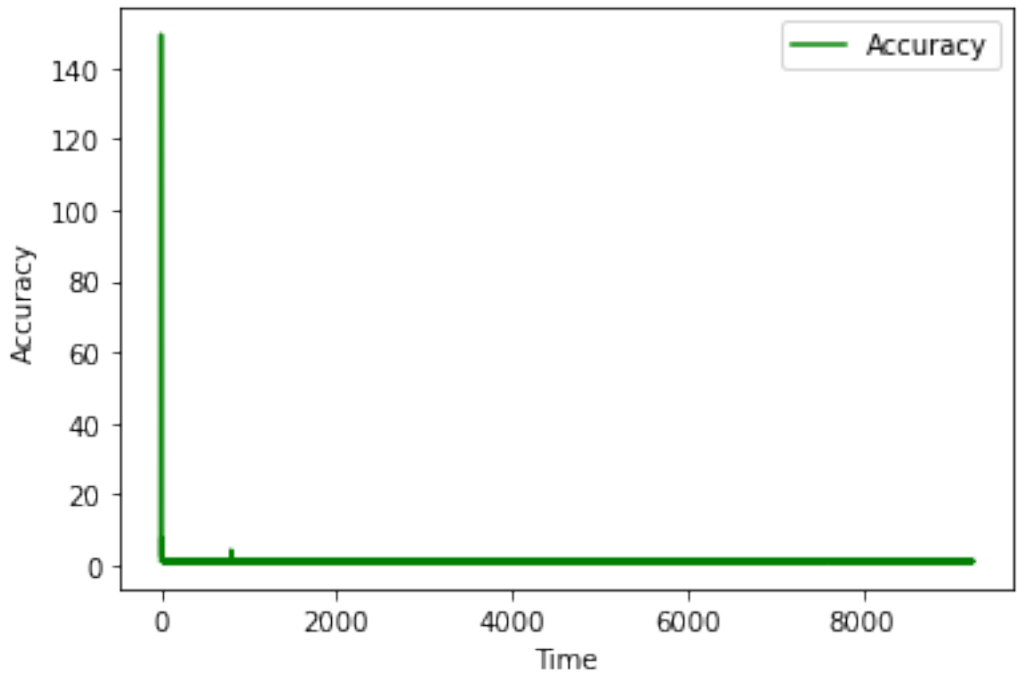

[9]:
```
with torch.no_grad():
    correct = 0
    total = 0
    for images, labels in test_loader:
        images = images.to(device)
        labels = labels.to(device)
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

    print("Test Accuracy",correct / total)

    correct = 0
    total = 0

    for images, labels in val_loader:
        images = images.to(device)
        labels = labels.to(device)
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
```

```
    print("Validation Accuracy",correct / total)

    correct = 0
    total = 0

    for images, labels in train_loader:
        images = images.to(device)
        labels = labels.to(device)
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

    print("Train Accuracy",correct / total)
```

```
Test Accuracy 0.485592820028342
Validation Accuracy 0.4899598393574297
Train Accuracy 0.4785690178872764
```

[10]:
```
def Show(out, title = ''):
    print(title)
    out = out.permute(1,0,2,3)
    grilla = torchvision.utils.make_grid(out)
    plt.imshow(transform.ToPILImage()(grilla), 'jet')
    plt.show()

def Show_Weight(out):
    grilla = torchvision.utils.make_grid(out)
    plt.imshow(transform.ToPILImage()(grilla), 'jet')
    plt.show()


with torch.no_grad():
    model.to('cpu')
    img, label = test_set[456]
    img = img.unsqueeze(0)
    out = model(img)
    print(out)
    print ((out == out.max()).nonzero())

    out = model.layer1[0](img)
    Show(out, 'layer 1: Convolution output')
    out = model.layer1[1](out)
    Show(out, 'layer 1: Activation function output')
```

```
    out = model.layer2[0](out)
    Show(out, 'layer 2: Convolution output')
    out = model.layer2[1](out)
    Show(out, 'layer 2: Activation function output')
    out = model.layer2[2](out)
    Show(out, 'layer 2: Max-Pooling')

    out = model.layer3[0](out)
    Show(out, 'layer 3: Convolution output')
    out = model.layer3[1](out)
    Show(out, 'layer 3: Activation function output')
    out = model.layer3[2](out)
    Show(out, 'layer 3: Max-Pooling')

    out = model.layer4[0](out)
    Show(out, 'layer 4: Convolution output')
    out = model.layer4[1](out)
    Show(out, 'layer 4: Activation function output')
    out = model.layer4[2](out)
    Show(out, 'layer 4: Max-Pooling')
```

```
tensor([[-0.2828,  0.1994,  0.6836, -1.3067]])
tensor([[0, 2]])
layer 1: Convolution output
```

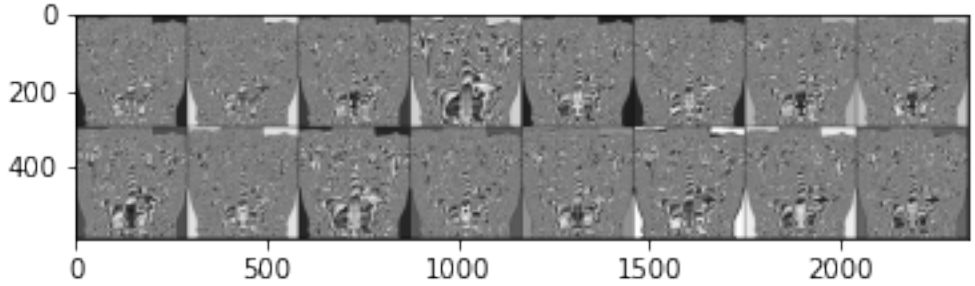


```
layer 1: Activation function output
```

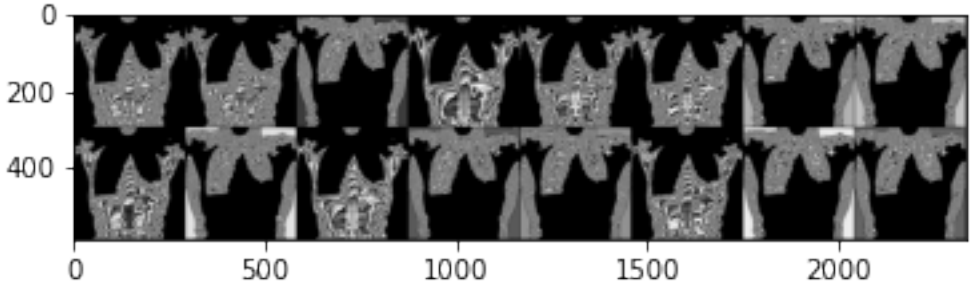

layer 2: Convolution output

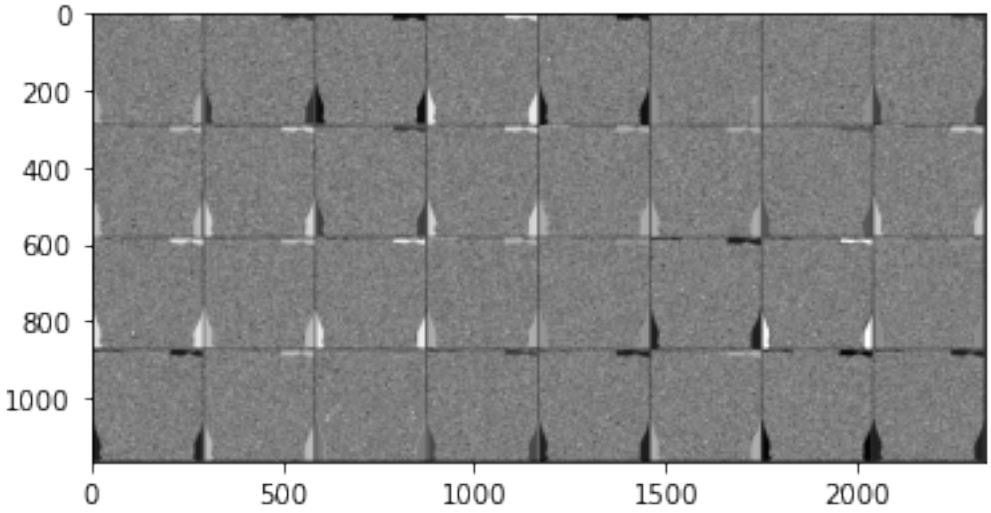

layer 2: Activation function output

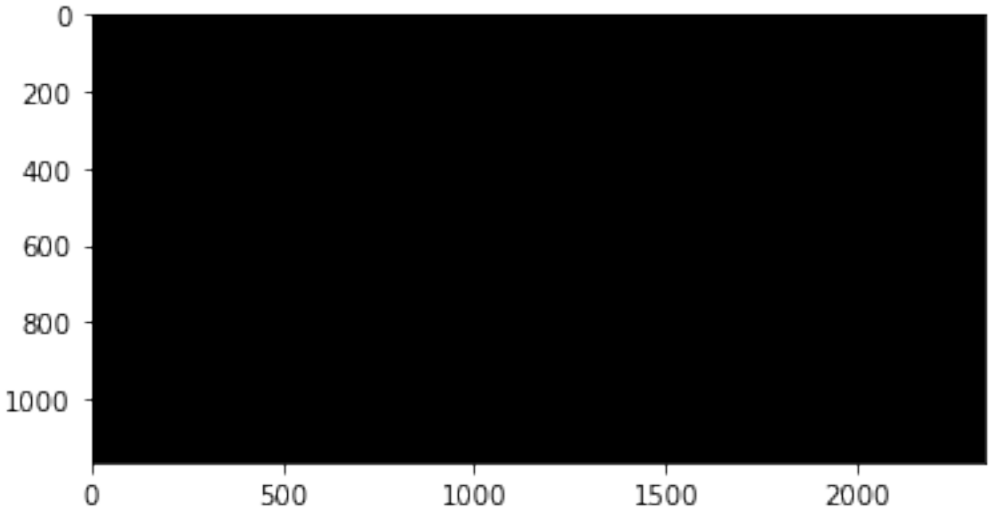

layer 2: Max-Pooling

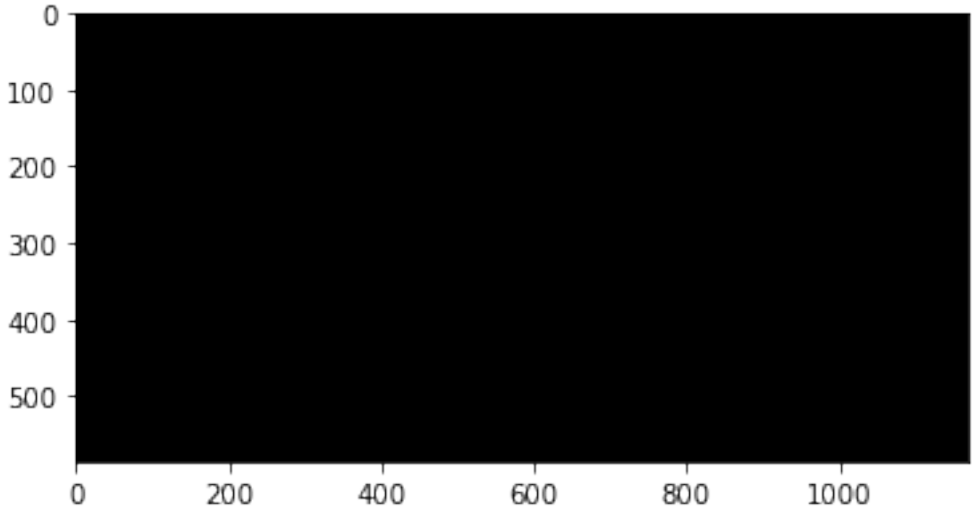

layer 3: Convolution output

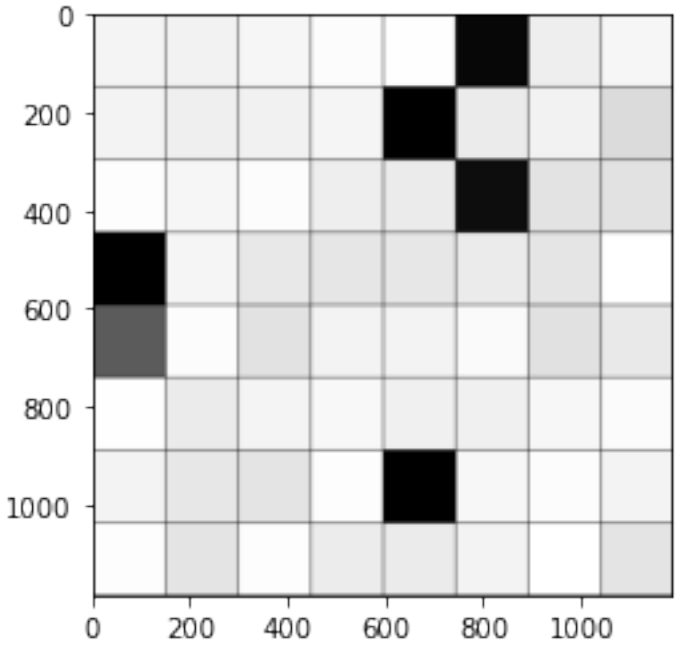

layer 3: Activation function output

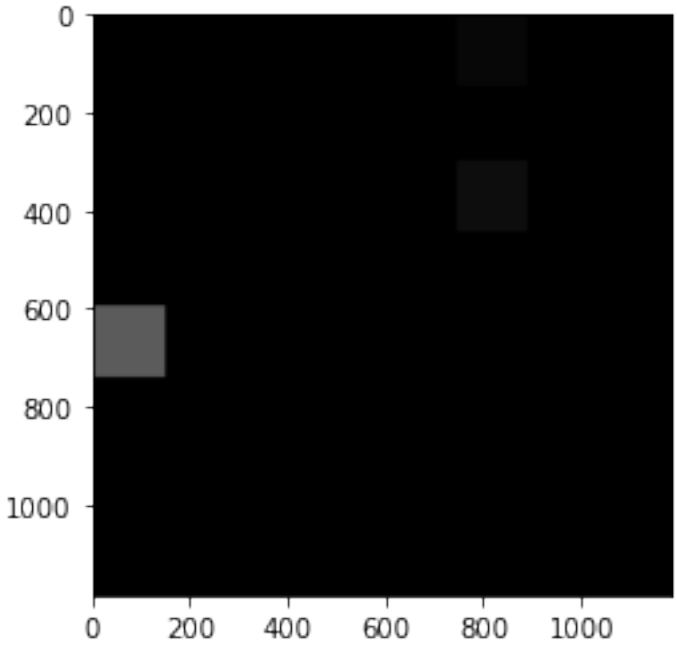

layer 3: Max-Pooling

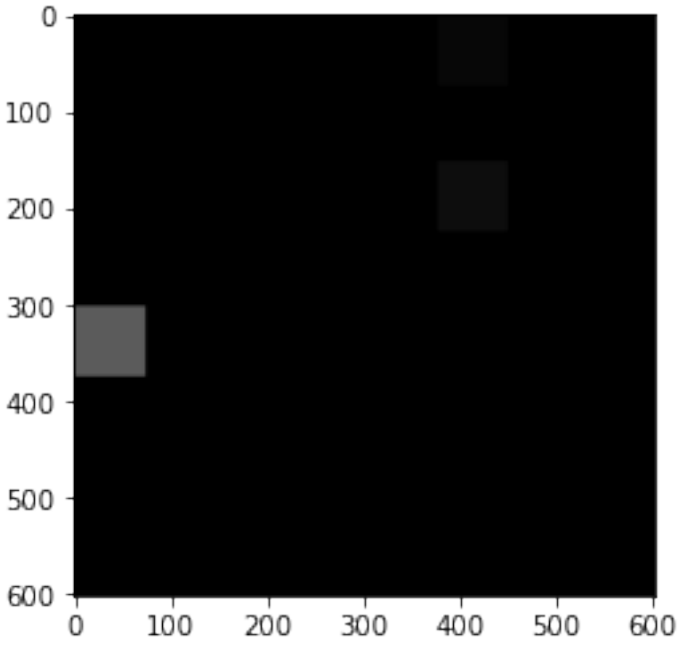

layer 4: Convolution output

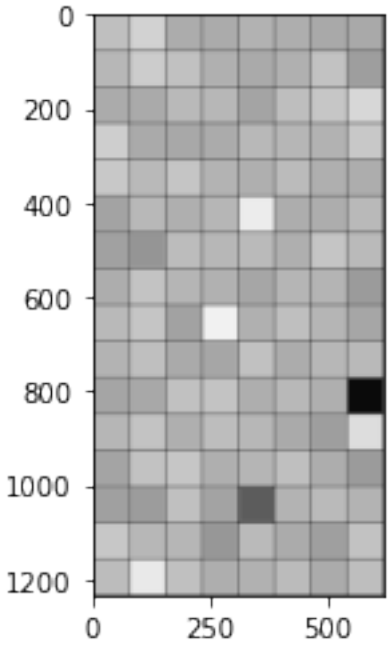

layer 4: Activation function output

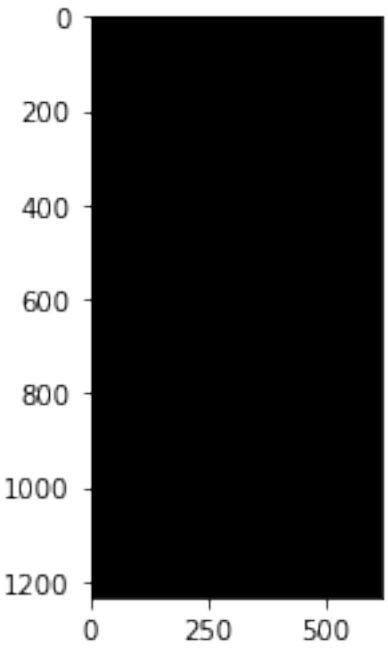

layer 4: Max-Pooling

[ ]: