

## Experimento 3

```
[1]: import torch
import torch.nn as nn
import torchvision
import torchvision.transforms as transform
import torch.nn.functional as F

import matplotlib.pyplot as plt
import numpy as np
import math

device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
#device='cpu'
print(device)

batch_size = 32

img_transform = transform.Compose([transform.ToTensor(), transform.Normalize((0.
→5,),(0.5,))])

data = torchvision.datasets.ImageFolder("C:/Users/matia/Desktop/AI_UTEC/
→proyecto5/fotos",transform=img_transform)

print(len(data))

train_set,test_set=torch.utils.data.random_split(data,[14815,6350],
→generator=torch.Generator().manual_seed(0))
val_set,test_set=torch.utils.data.random_split(test_set,[4233,2117],
→generator=torch.Generator().manual_seed(0))

img, _ = train_set[0]
print(img.shape)

train_loader = torch.utils.data.DataLoader(dataset=train_set,
→batch_size=batch_size, shuffle=True)
test_loader = torch.utils.data.DataLoader(dataset=test_set,
→batch_size=batch_size, shuffle=False)
```

```
val_loader = torch.utils.data.DataLoader(dataset=val_set, batch_size=batch_size,
→shuffle=False)
```

```
cuda:0
21165
torch.Size([3, 299, 299])
```


```
[2]: def show_img(img):
      plt.imshow(img.numpy()[0], cmap='gray')
```

```
[3]: print(len(train_set))
      print(len(test_set))
      print(len(val_set))
```

```
14815
2117
4233
```

```
[4]: img, label = train_set[999]
      print(label)
      show_img(img)
```

```
3
```



```
[5]: #hyperparametros
      num_classes = 4
      learning_rate = 0.01
```

```

num_epochs = 20

class CNN(nn.Module):
    def __init__(self, num_classes=4):
        super(CNN, self).__init__()
        self.layer1 = nn.Sequential(
            nn.Conv2d(in_channels=3, out_channels=16, kernel_size=10,
→stride=1,padding=0),
            nn.ReLU(),nn.BatchNorm2d(16))
        self.layer2 = nn.Sequential(
            nn.Conv2d(in_channels=16,out_channels=32, kernel_size=6, stride=1,
→padding=2),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2))
        self.layer3 = nn.Sequential(
            nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3, stride=1,
→padding=2),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2))
        self.layer4 = nn.Sequential(
            nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3, stride=1,
→padding=2),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2))
        self.fc = nn.Linear(37*37*128, num_classes)

    def forward(self, x):
        out = self.layer1(x)
        out = self.layer2(out)
        out = self.layer3(out)
        out = self.layer4(out)
        out = out.reshape(out.size(0), -1)
        out = self.fc(out)
        return out

```

```

[6]: model = CNN(num_classes).to(device)
loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr = learning_rate)
#loss_train = train(model, optimizer, loss_fn, num_epochs)
#test(model)

print([ e.shape for e in model.fc.parameters()])

model.fc.weight

```

```

[torch.Size([4, 175232]), torch.Size([4])]

```

```
[6]: Parameter containing:
      tensor([[ -9.4709e-04,  2.1630e-03, -1.2115e-03, ..., -1.5297e-03,
                -1.3633e-03,  1.4940e-03],
              [ 1.4682e-03,  2.1715e-03, -1.7284e-04, ..., -2.1634e-03,
                -2.1812e-04,  1.9144e-03],
              [-1.7999e-03,  8.1053e-04, -9.3822e-05, ..., -2.1619e-03,
                -1.8192e-03, -2.1854e-03],
              [ 1.0879e-03,  6.7307e-04, -1.9577e-03, ...,  1.1216e-03,
                2.2792e-03, -1.4664e-03]], device='cuda:0', requires_grad=True)
```

```
[7]: def train(model, optimizer, loss_fn, num_epochs):
      loss_vals = []
      running_loss = 0.0
      # train the model
      total_step = len(train_loader)

      list_loss = []
      list_time = []
      j = 0

      for epoch in range(num_epochs):
          for i, (images, labels) in enumerate(train_loader):
              images = images.to(device)
              labels = labels.to(device)
              # forward
              output = model(images)
              loss = loss_fn(output, labels)
              # change the params
              optimizer.zero_grad()
              loss.backward()
              optimizer.step()

              list_loss.append(loss.item())
              list_time.append(j)
              j += 1
              # print(i, end=" ")
              if (i+1) % 100 == 0:
                  # print()
                  print ('Epoch [{}/{}], Step [{}/{}], Loss: {:.4f}' .
→ format(epoch+1, num_epochs, i+1, total_step, loss.item()))

      print('Finished Training Trainset')
      return list_loss, list_time
```


```
[8]: list_loss, list_time = train(model, optimizer, loss_fn, num_epochs)

      plt.plot(list_time, list_loss, color="green", label="Accuracy")
```

```
plt.legend()  
plt.xlabel("Time")  
plt.ylabel("Accuracy")  
plt.show()
```

```
Epoch [1/20], Step [100/463], Loss: 1.2155  
Epoch [1/20], Step [200/463], Loss: 1.2081  
Epoch [1/20], Step [300/463], Loss: 1.1235  
Epoch [1/20], Step [400/463], Loss: 1.0604  
Epoch [2/20], Step [100/463], Loss: 1.2279  
Epoch [2/20], Step [200/463], Loss: 1.1592  
Epoch [2/20], Step [300/463], Loss: 1.2901  
Epoch [2/20], Step [400/463], Loss: 1.1010  
Epoch [3/20], Step [100/463], Loss: 1.1965  
Epoch [3/20], Step [200/463], Loss: 1.1855  
Epoch [3/20], Step [300/463], Loss: 1.1324  
Epoch [3/20], Step [400/463], Loss: 1.0923  
Epoch [4/20], Step [100/463], Loss: 1.1264  
Epoch [4/20], Step [200/463], Loss: 1.4292  
Epoch [4/20], Step [300/463], Loss: 1.2406  
Epoch [4/20], Step [400/463], Loss: 1.1045  
Epoch [5/20], Step [100/463], Loss: 1.1081  
Epoch [5/20], Step [200/463], Loss: 1.1022  
Epoch [5/20], Step [300/463], Loss: 1.0910  
Epoch [5/20], Step [400/463], Loss: 1.0902  
Epoch [6/20], Step [100/463], Loss: 1.0136  
Epoch [6/20], Step [200/463], Loss: 1.0852  
Epoch [6/20], Step [300/463], Loss: 0.9771  
Epoch [6/20], Step [400/463], Loss: 1.1041  
Epoch [7/20], Step [100/463], Loss: 1.0270  
Epoch [7/20], Step [200/463], Loss: 1.0492  
Epoch [7/20], Step [300/463], Loss: 1.0466  
Epoch [7/20], Step [400/463], Loss: 1.0625  
Epoch [8/20], Step [100/463], Loss: 0.9880  
Epoch [8/20], Step [200/463], Loss: 0.8552  
Epoch [8/20], Step [300/463], Loss: 1.4600  
Epoch [8/20], Step [400/463], Loss: 1.0644  
Epoch [9/20], Step [100/463], Loss: 1.0430  
Epoch [9/20], Step [200/463], Loss: 0.7856  
Epoch [9/20], Step [300/463], Loss: 0.8781  
Epoch [9/20], Step [400/463], Loss: 1.1859  
Epoch [10/20], Step [100/463], Loss: 0.9191  
Epoch [10/20], Step [200/463], Loss: 0.9157  
Epoch [10/20], Step [300/463], Loss: 0.9650  
Epoch [10/20], Step [400/463], Loss: 0.7295  
Epoch [11/20], Step [100/463], Loss: 0.9834  
Epoch [11/20], Step [200/463], Loss: 0.8449  
Epoch [11/20], Step [300/463], Loss: 0.9846
```

Epoch [11/20], Step [400/463], Loss: 0.9690  
Epoch [12/20], Step [100/463], Loss: 1.0576  
Epoch [12/20], Step [200/463], Loss: 1.0967  
Epoch [12/20], Step [300/463], Loss: 0.9086  
Epoch [12/20], Step [400/463], Loss: 0.9322  
Epoch [13/20], Step [100/463], Loss: 1.0330  
Epoch [13/20], Step [200/463], Loss: 1.1576  
Epoch [13/20], Step [300/463], Loss: 0.7888  
Epoch [13/20], Step [400/463], Loss: 0.9810  
Epoch [14/20], Step [100/463], Loss: 1.1704  
Epoch [14/20], Step [200/463], Loss: 0.8800  
Epoch [14/20], Step [300/463], Loss: 1.0530  
Epoch [14/20], Step [400/463], Loss: 0.9399  
Epoch [15/20], Step [100/463], Loss: 0.8664  
Epoch [15/20], Step [200/463], Loss: 0.9916  
Epoch [15/20], Step [300/463], Loss: 0.6895  
Epoch [15/20], Step [400/463], Loss: 1.0541  
Epoch [16/20], Step [100/463], Loss: 0.9471  
Epoch [16/20], Step [200/463], Loss: 1.0457  
Epoch [16/20], Step [300/463], Loss: 0.8010  
Epoch [16/20], Step [400/463], Loss: 1.1985  
Epoch [17/20], Step [100/463], Loss: 1.0387  
Epoch [17/20], Step [200/463], Loss: 1.0713  
Epoch [17/20], Step [300/463], Loss: 0.9556  
Epoch [17/20], Step [400/463], Loss: 0.7234  
Epoch [18/20], Step [100/463], Loss: 1.0660  
Epoch [18/20], Step [200/463], Loss: 0.8599  
Epoch [18/20], Step [300/463], Loss: 1.1548  
Epoch [18/20], Step [400/463], Loss: 0.9046  
Epoch [19/20], Step [100/463], Loss: 0.6636  
Epoch [19/20], Step [200/463], Loss: 1.0473  
Epoch [19/20], Step [300/463], Loss: 0.8127  
Epoch [19/20], Step [400/463], Loss: 1.1208  
Epoch [20/20], Step [100/463], Loss: 0.8305  
Epoch [20/20], Step [200/463], Loss: 1.0279  
Epoch [20/20], Step [300/463], Loss: 0.6081  
Epoch [20/20], Step [400/463], Loss: 0.8848  
Finished Training Trainset



```
[9]: with torch.no_grad():
    correct = 0
    total = 0
    for images, labels in test_loader:
        images = images.to(device)
        labels = labels.to(device)
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

    print("Test Accuracy", correct / total)

    correct = 0
    total = 0

    for images, labels in val_loader:
        images = images.to(device)
        labels = labels.to(device)
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
```

```

print("Validation Accuracy",correct / total)

correct = 0
total = 0

for images, labels in train_loader:
    images = images.to(device)
    labels = labels.to(device)
    outputs = model(images)
    _, predicted = torch.max(outputs.data, 1)
    total += labels.size(0)
    correct += (predicted == labels).sum().item()

print("Train Accuracy",correct / total)

```

Test Accuracy 0.596126594237128  
Validation Accuracy 0.5931963146704465  
Train Accuracy 0.5801552480593992

```

[10]: def Show(out, title = ''):
    print(title)
    out = out.permute(1,0,2,3)
    grilla = torchvision.utils.make_grid(out)
    plt.imshow(transform.ToPILImage()(grilla), 'jet')
    plt.show()

def Show_Weight(out):
    grilla = torchvision.utils.make_grid(out)
    plt.imshow(transform.ToPILImage()(grilla), 'jet')
    plt.show()

with torch.no_grad():
    model.to('cpu')
    img, label = test_set[456]
    img = img.unsqueeze(0)
    out = model(img)
    print(out)
    print ((out == out.max()).nonzero())

    out = model.layer1[0](img)
    Show(out, 'layer 1: Convolution output')
    out = model.layer1[1](out)
    Show(out, 'layer 1: Activation function output')

```



```

out = model.layer2[0](out)
Show(out, 'layer 2: Convolution output')
out = model.layer2[1](out)
Show(out, 'layer 2: Activation function output')
out = model.layer2[2](out)
Show(out, 'layer 2: Max-Pooling')

out = model.layer3[0](out)
Show(out, 'layer 3: Convolution output')
out = model.layer3[1](out)
Show(out, 'layer 3: Activation function output')
out = model.layer3[2](out)
Show(out, 'layer 3: Max-Pooling')


out = model.layer4[0](out)
Show(out, 'layer 4: Convolution output')
out = model.layer4[1](out)
Show(out, 'layer 4: Activation function output')
out = model.layer4[2](out)
Show(out, 'layer 4: Max-Pooling')

```


```
tensor([[ -0.3140, -1.7436,  0.6980, -3.1916]])
```

```
tensor([[0, 2]])
```


```
layer 1: Convolution output
```




```
layer 1: Activation function output
```




layer 2: Convolution output




layer 2: Activation function output




layer 2: Max-Pooling




layer 3: Convolution output




layer 3: Activation function output



layer 3: Max-Pooling




layer 4: Convolution output



layer 4: Activation function output



layer 4: Max-Pooling



[ ]:

