

## Experimento 4

```
[1]: import torch
import torch.nn as nn
import torchvision
import torchvision.transforms as transform
import torch.nn.functional as F

import matplotlib.pyplot as plt
import numpy as np
import math

device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
#device='cpu'
print(device)

batch_size = 32

img_transform = transform.Compose([transform.ToTensor(), transform.Normalize((0.
→5,),(0.5,))])

data = torchvision.datasets.ImageFolder("C:/Users/matia/Desktop/AI_UTEC/
→proyecto5/fotos",transform=img_transform)

print(len(data))

train_set,test_set=torch.utils.data.random_split(data,[14815,6350],
→generator=torch.Generator().manual_seed(0))
val_set,test_set=torch.utils.data.random_split(test_set,[4233,2117],
→generator=torch.Generator().manual_seed(0))

img, _ = train_set[0]
print(img.shape)

train_loader = torch.utils.data.DataLoader(dataset=train_set,
→batch_size=batch_size, shuffle=True)
test_loader = torch.utils.data.DataLoader(dataset=test_set,
→batch_size=batch_size, shuffle=False)
```

```
val_loader = torch.utils.data.DataLoader(dataset=val_set, batch_size=batch_size,
→shuffle=False)
```

```
cuda:0
21165
torch.Size([3, 299, 299])
```


```
[2]: def show_img(img):
      plt.imshow(img.numpy()[0], cmap='gray')
```

```
[3]: print(len(train_set))
      print(len(test_set))
      print(len(val_set))
```

```
14815
2117
4233
```

```
[4]: img, label = train_set[999]
      print(label)
      show_img(img)
```

```
3
```



```
[5]: #hyperparametros
      num_classes = 4
      learning_rate = 0.01
```

```

num_epochs = 20

class CNN(nn.Module):
    def __init__(self, num_classes=4):
        super(CNN, self).__init__()
        self.layer1 = nn.Sequential(
            nn.Conv2d(in_channels=3, out_channels=16, kernel_size=10,
→stride=1,padding=0),
            nn.ReLU(),nn.BatchNorm2d(16))
        self.layer2 = nn.Sequential(
            nn.Conv2d(in_channels=16,out_channels=32, kernel_size=6, stride=1,
→padding=2),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2),nn.BatchNorm2d(32))
        self.layer3 = nn.Sequential(
            nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3, stride=1,
→padding=2),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2),nn.BatchNorm2d(64))
        self.layer4 = nn.Sequential(
            nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3, stride=1,
→padding=2),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2),nn.BatchNorm2d(128))
        self.fc = nn.Linear(37*37*128, num_classes)

    def forward(self, x):
        out = self.layer1(x)
        out = self.layer2(out)
        out = self.layer3(out)
        out = self.layer4(out)
        out = out.reshape(out.size(0), -1)
        out = self.fc(out)
        return out

```

```

[6]: model = CNN(num_classes).to(device)
loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr = learning_rate)
#loss_train = train(model, optimizer, loss_fn, num_epochs)
#test(model)

print([ e.shape for e in model.fc.parameters()])

model.fc.weight

```

```

[torch.Size([4, 175232]), torch.Size([4])]

```

```
[6]: Parameter containing:
      tensor([[ -0.0010,  0.0014,  0.0012, ..., -0.0020, -0.0016, -0.0003],
              [ 0.0003, -0.0015,  0.0016, ..., -0.0008, -0.0013, -0.0018],
              [-0.0012, -0.0020,  0.0020, ..., -0.0013, -0.0021,  0.0021],
              [-0.0023, -0.0016, -0.0015, ..., -0.0010,  0.0003,  0.0005]],
              device='cuda:0', requires_grad=True)
```

```
[7]: def train(model, optimizer, loss_fn, num_epochs):
      loss_vals = []
      running_loss = 0.0
      # train the model
      total_step = len(train_loader)

      list_loss= []
      list_time = []
      j=0

      for epoch in range(num_epochs):
          for i, (images, labels) in enumerate(train_loader):
              images = images.to(device)
              labels = labels.to(device)
              # forward
              output = model(images)
              loss = loss_fn(output, labels)
              # change the params
              optimizer.zero_grad()
              loss.backward()
              optimizer.step()

              list_loss.append(loss.item())
              list_time.append(j)
              j+=1
              #print(i,end=",")
              if (i+1) % 100 == 0:
                  #print()
                  print ('Epoch [{}/{}], Step [{}/{}], Loss: {:.4f}' .
→format(epoch+1, num_epochs, i+1, total_step, loss.item()))

      print('Finished Training Trainset')
      return list_loss,list_time
```


```
[8]: list_loss,list_time=train(model,optimizer,loss_fn,num_epochs)

plt.plot(list_time,list_loss,color="green", label="Accuracy")
plt.legend()
plt.xlabel("Time")
plt.ylabel("Accuracy")
```

```
plt.show()
```

```
Epoch [1/20], Step [100/463], Loss: 26.5680
Epoch [1/20], Step [200/463], Loss: 17.9783
Epoch [1/20], Step [300/463], Loss: 4.2819
Epoch [1/20], Step [400/463], Loss: 0.7125
Epoch [2/20], Step [100/463], Loss: 0.4105
Epoch [2/20], Step [200/463], Loss: 0.6548
Epoch [2/20], Step [300/463], Loss: 0.3909
Epoch [2/20], Step [400/463], Loss: 0.8094
Epoch [3/20], Step [100/463], Loss: 0.3243
Epoch [3/20], Step [200/463], Loss: 0.3466
Epoch [3/20], Step [300/463], Loss: 0.6585
Epoch [3/20], Step [400/463], Loss: 1.0781
Epoch [4/20], Step [100/463], Loss: 0.3112
Epoch [4/20], Step [200/463], Loss: 0.4940
Epoch [4/20], Step [300/463], Loss: 0.5900
Epoch [4/20], Step [400/463], Loss: 0.1904
Epoch [5/20], Step [100/463], Loss: 0.5739
Epoch [5/20], Step [200/463], Loss: 0.6155
Epoch [5/20], Step [300/463], Loss: 0.3001
Epoch [5/20], Step [400/463], Loss: 0.3333
Epoch [6/20], Step [100/463], Loss: 0.5945
Epoch [6/20], Step [200/463], Loss: 0.2525
Epoch [6/20], Step [300/463], Loss: 0.8553
Epoch [6/20], Step [400/463], Loss: 0.6627
Epoch [7/20], Step [100/463], Loss: 0.4400
Epoch [7/20], Step [200/463], Loss: 0.6392
Epoch [7/20], Step [300/463], Loss: 1.1934
Epoch [7/20], Step [400/463], Loss: 0.1875
Epoch [8/20], Step [100/463], Loss: 0.7594
Epoch [8/20], Step [200/463], Loss: 0.5291
Epoch [8/20], Step [300/463], Loss: 0.7889
Epoch [8/20], Step [400/463], Loss: 0.6325
Epoch [9/20], Step [100/463], Loss: 0.1640
Epoch [9/20], Step [200/463], Loss: 1.0866
Epoch [9/20], Step [300/463], Loss: 0.9030
Epoch [9/20], Step [400/463], Loss: 1.5918
Epoch [10/20], Step [100/463], Loss: 0.7051
Epoch [10/20], Step [200/463], Loss: 2.0270
Epoch [10/20], Step [300/463], Loss: 0.2571
Epoch [10/20], Step [400/463], Loss: 1.0388
Epoch [11/20], Step [100/463], Loss: 0.2761
Epoch [11/20], Step [200/463], Loss: 0.6677
Epoch [11/20], Step [300/463], Loss: 0.7748
Epoch [11/20], Step [400/463], Loss: 0.7447
Epoch [12/20], Step [100/463], Loss: 1.9397
Epoch [12/20], Step [200/463], Loss: 1.0744
```

Epoch [12/20], Step [300/463], Loss: 0.2082  
Epoch [12/20], Step [400/463], Loss: 2.4821  
Epoch [13/20], Step [100/463], Loss: 1.6145  
Epoch [13/20], Step [200/463], Loss: 1.7571  
Epoch [13/20], Step [300/463], Loss: 1.2336  
Epoch [13/20], Step [400/463], Loss: 1.3819  
Epoch [14/20], Step [100/463], Loss: 0.5558  
Epoch [14/20], Step [200/463], Loss: 1.9162  
Epoch [14/20], Step [300/463], Loss: 0.7344  
Epoch [14/20], Step [400/463], Loss: 0.0889  
Epoch [15/20], Step [100/463], Loss: 0.6145  
Epoch [15/20], Step [200/463], Loss: 0.2542  
Epoch [15/20], Step [300/463], Loss: 0.3374  
Epoch [15/20], Step [400/463], Loss: 1.1852  
Epoch [16/20], Step [100/463], Loss: 0.4219  
Epoch [16/20], Step [200/463], Loss: 0.5924  
Epoch [16/20], Step [300/463], Loss: 0.4941  
Epoch [16/20], Step [400/463], Loss: 0.7897  
Epoch [17/20], Step [100/463], Loss: 1.4301  
Epoch [17/20], Step [200/463], Loss: 1.1623  
Epoch [17/20], Step [300/463], Loss: 4.7026  
Epoch [17/20], Step [400/463], Loss: 0.3348  
Epoch [18/20], Step [100/463], Loss: 0.2741  
Epoch [18/20], Step [200/463], Loss: 0.6788  
Epoch [18/20], Step [300/463], Loss: 0.8928  
Epoch [18/20], Step [400/463], Loss: 0.2067  
Epoch [19/20], Step [100/463], Loss: 0.1447  
Epoch [19/20], Step [200/463], Loss: 0.0273  
Epoch [19/20], Step [300/463], Loss: 0.6088  
Epoch [19/20], Step [400/463], Loss: 1.0564  
Epoch [20/20], Step [100/463], Loss: 1.4888  
Epoch [20/20], Step [200/463], Loss: 1.6213  
Epoch [20/20], Step [300/463], Loss: 0.9583  
Epoch [20/20], Step [400/463], Loss: 0.8872  
Finished Training Trainset



```
[9]: with torch.no_grad():
    correct = 0
    total = 0
    for images, labels in test_loader:
        images = images.to(device)
        labels = labels.to(device)
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

    print("Test Accuracy", correct / total)

    correct = 0
    total = 0

    for images, labels in val_loader:
        images = images.to(device)
        labels = labels.to(device)
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
```

```

print("Validation Accuracy",correct / total)

correct = 0
total = 0

for images, labels in train_loader:
    images = images.to(device)
    labels = labels.to(device)
    outputs = model(images)
    _, predicted = torch.max(outputs.data, 1)
    total += labels.size(0)
    correct += (predicted == labels).sum().item()

print("Train Accuracy",correct / total)

```

Test Accuracy 0.8375059045819556  
Validation Accuracy 0.8230569336168202  
Train Accuracy 0.9561930475869052

```

[16]: def Show(out, title = ''):
    print(title)
    out = out.permute(1,0,2,3)
    grilla = torchvision.utils.make_grid(out)
    plt.imshow(transform.ToPILImage()(grilla), 'jet')
    plt.show()

def Show_Weight(out):
    grilla = torchvision.utils.make_grid(out)
    plt.imshow(transform.ToPILImage()(grilla), 'jet')
    plt.show()

with torch.no_grad():
    model.to('cpu')
    img, label = test_set[456]
    img = img.unsqueeze(0)
    out = model(img)
    print(out)
    print ((out == out.max()).nonzero())

    out = model.layer1[0](img)
    Show(out, 'layer 1: Convolution output')
    out = model.layer1[1](out)
    Show(out, 'layer 1: Activation function output')

```



```

out = model.layer2[0](out)
Show(out, 'layer 2: Convolution output')
out = model.layer2[1](out)
Show(out, 'layer 2: Activation function output')
out = model.layer2[2](out)
Show(out, 'layer 2: Max-Pooling')

out = model.layer3[0](out)
Show(out, 'layer 3: Convolution output')
out = model.layer3[1](out)
Show(out, 'layer 3: Activation function output')
out = model.layer3[2](out)
Show(out, 'layer 3: Max-Pooling')


out = model.layer4[0](out)
Show(out, 'layer 4: Convolution output')
out = model.layer4[1](out)
Show(out, 'layer 4: Activation function output')
out = model.layer4[2](out)
Show(out, 'layer 4: Max-Pooling')

```


```
tensor([[ -59.4509,  22.0876,  24.1623, -76.1220]])
```

```
tensor([[0, 2]])
```


```
layer 1: Convolution output
```




```
layer 1: Activation function output
```




layer 2: Convolution output




layer 2: Activation function output




layer 2: Max-Pooling




layer 3: Convolution output




layer 3: Activation function output




layer 3: Max-Pooling




layer 4: Convolution output



layer 4: Activation function output



layer 4: Max-Pooling



[ ]: